


List of deliverables

Project: Recruit-AI



Problem framing document



USER PERSONA: Sarah Mitchell

"I want to hire the best candidates for job, but I am spending all my time just trying to read the rest."



AT A GLANCE

Attribute	Details
Role	Solo Talent Acquisition Manager
Company	Fast-growing 50-person tech startup
Location	Remote / Hybrid (High-volume hiring)
KPI Goal	Double company headcount in 6 months



TRAITS & LITERACY

- **Tech Savviness:**  80% (*Power user of SaaS, but not a developer*)
- **Daily Workload:**  100% (*Severely overwhelmed*)
- **Automation Use:**  20% (*Highly manual spreadsheet processes*)



THE FRICTION (Pain Points)

- 🗑️ **The Volume Crush:** Receives **100-500+ unstructured PDFs** per open role.
- ⌚ **The Time Drain:** Spends **20+ hours/week** manually scanning (averaging just 6 seconds per resume).
- 🧠 **Decision Fatigue:** By the 100th resume, her eyes glaze over, risking the accidental rejection of top-tier talent.
- 💰 **The Cost of Delay:** Top engineers accept competing offers within 14 days because her manual screening process takes 3+ weeks.



THE FIX (Goals & Needs)

- 🤖 **Agentic AI Assistant:** Needs a tool to autonomously read, parse, and score incoming PDFs against specific Job Descriptions.

- 📁 **Instant Shortlisting:** Wants to spend her limited time *only* interviewing the top 10% of matches.
- 🕒 **Speed to Hire:** Must reduce average time-to-hire from **32 days to under 15 days**.

🕒 A DAY IN THE LIFE (The Bottleneck)

- **09:00 AM** ☕ Opens inbox. Sees 150 new applications. Anxiety spikes.
- **10:30 AM** 📄 Opens PDFs one by one. Mentally compares unstructured text to rigid Job Descriptions.
- **02:00 PM** 🧩 Tries to organize "Maybe" candidates using messy spreadsheets and sticky notes.
- **04:30 PM** ❌ Discovers her top pick accepted another offer because she took too long to schedule the interview.

The Problem Statement: The High-Volume Hiring Bottleneck

1. The Volume and Capacity Paradox Mid-sized tech companies are experiencing a critical operational failure in the "screening and scheduling" phase of talent acquisition. When sourcing for competitive IT and engineering roles, a single job posting routinely attracts 250 to 500+ applications. For lean HR teams—often operating with a solo recruiter like Sarah—this massive influx of unstructured applicant data makes it physically impossible to evaluate candidates thoroughly while maintaining the speed required to secure top talent.

2. The Flaws of Current Screening Methods Existing solutions fail to bridge this gap. Traditional Applicant Tracking Systems (ATS) act merely as passive digital filing cabinets that rely on rigid, keyword-matching filters, frequently rejecting highly qualified candidates due to simple formatting errors or unconventional career paths. Consequently, recruiters are forced to manually open and skim hundreds of PDF resumes, spending an average of just 6 to 8 seconds per document. This highly repetitive manual process consumes upwards of 23 hours per open role and induces severe decision fatigue, leading to inconsistent evaluations, unconscious bias, and the accidental rejection of elite candidates.

3. The Business and Financial Impact This administrative friction creates a crippling domino effect on the business. Because the initial screening process is heavily manual, the overall average time-to-hire stretches to an agonizing 32 days. In a hyper-competitive market where top-tier talent accepts competing offers within the first 14 days, this delay is fatal. Companies not only experience massive candidate drop-off rates and damaged employer branding, but they also incur severe financial losses—costing the business roughly ₹3-6 lakh per month in lost productivity per vacant role. Ultimately, hiring managers are forced to make desperate, lower-quality hires simply to fill a seat.

Recruit-AI Market Research Summary

The global HR Tech market is valued at \$43.6B in 2025 (8% CAGR to \$80.2B by 2033), with AI recruitment growing at 6.8% CAGR to \$1.1B by 2032. Asia Pacific holds 23% share (\$27.7B in 2024), while India's HR Tech market (\$1.12B in 2024) doubles to \$2.3B by 2033 (7.56% CAGR); recruitment software hits \$81.6M in 2025 (4.52% CAGR to \$122.9M by 2034). Key Indian pain: 100-500+ applicants per IT/sales role (up 48% in downturns), 32-day avg time-to-hire, talent shortages for 80% employers.

Core Friction: Recruiters waste 23 hrs/role on manual screening (4-5 hrs initial scans at 6-8 sec/resume for 500 apps), costing ₹3-6L/month per vacancy. Decision fatigue causes rushed hires/turnover; 60% top talent drops after 3 weeks globally (14 days India).

Competitors (Traditional ATS lack agentic AI, rely on manual filters/integrations):

Category	Key Players	Focus & Gaps	Region
Traditional ATS	Greenhouse, Workable, Lever, iCIMS, JazzHR, SmartRecruiters, Recrutee, ClearCompany skima +2	Tracking/onboarding; weak autonomous screening.	Global
Global AI/Agentic	Paradox.ai (conversational scheduling), Eightfold.ai (skills intel/mobility), Braintrust, SorsX, hireEZ, Phenom, SeekOut, Gem, HireVue, Beamery eightfold +6	End-to-end automation, sourcing/outreach.	Worldwide
Asia/Global AI	Apna.co (AI shortlisting, 50% time cut), Joveo (noise reduction) joveo +1	Job matching for high-volume.	India/Asia
India-Specific	Skillate (resume parsing), Belong.co (social matching, 15-20 hrs/week saved) indbiz	Volume handling via algorithms.	India

Pitch deck

Recruit-AI

Autonomous Candidate Screening Agent

Solving the "**High-Volume Hiring**" operational bottleneck — with agentic AI that works at the speed of your best competitors.



Made with GRAMMA

The Human Cost of Manual Screening

Meet Sarah

Solo recruiter at a 50-person startup. For every open role, she faces an avalanche of **500+ unstructured PDF resumes** — each one demanding her full attention.

- Spends **20+ hours weekly** on manual resume scanning
- Averages just **6 seconds per resume** under volume pressure
- Top candidates slip through — or accept offers elsewhere first

The Real Consequence

This manual crush isn't just inefficient — it's a **competitive liability**. While Sarah reviews page 47 of 500, a faster competitor has already extended an offer to your best candidate.


- ❑ Speed-to-offer is the new talent advantage. Manual screening is the bottleneck standing in the way.



Made with GRAMMA

The \$1.1 Billion AI Gap

MARKET OPPORTUNITY



The AI recruitment market is projected to reach **\$1.1B by 2032**, driven by surging applicant volumes and demand for intelligent automation.

Why Act Now?

Applicant Volume Doubling

High-volume markets like India are seeing applicant counts double YoY — legacy tools can't keep pace.

ATS = Passive Filing Cabinets

Legacy Applicant Tracking Systems store resumes — they don't evaluate them. The gap between storage and intelligence is wide open.

First-Mover Advantage

Teams that automate screening now will lock in a structural hiring speed advantage over competitors still reviewing PDFs manually.



An Autonomous AI Junior Recruiter

THE SOLUTION

Legacy ATS: Rigid & Passive

- Keyword matching only — misses contextual fit
- No scoring, no ranking, no synthesis
- Requires recruiter to manually read every resume


Recruit-AI: Contextual & Active

- LLM reads resumes contextually — like a human senior recruiter
- Autonomously parses, scores against the JD, and summarizes
- Delivers ranked candidates with justified scores — instantly

Made with **GRAMMAR**

The User Experience: From Weeks to Seconds

FRONTEND PROTOTYPE




Step 1 — Upload

Recruiter drags a PDF resume into a clean dropzone built in **Lovable (React)**. No formatting required. Any resume, any structure.



Step 2 — Process

A **3-second processing state** displays while the backend workflow parses, evaluates, and scores the candidate against the active Job Description.



Step 3 — Review

A ranked table appears instantly: **Match Score (e.g., 88%)**, a status badge ("INTERVIEW"), and an expandable AI-written candidate summary.

Made with **GRAMMAR**

Under the Hood: Low-Code, High-Impact

TECHNICAL ARCHITECTURE



Webhook

Parse PDF

Call Gemini

Return
JSON

Orchestrated entirely via **n8n**, this low-code workflow eliminates custom infrastructure overhead while delivering enterprise-grade AI evaluation. Each step is independently testable, auditable, and replaceable — giving the engineering team full visibility and control without vendor lock-in.

Made with **GRAMMA**

Beyond Keyword Matching

AGENTIC LOGIC

The Prompt Engineering Strategy

We don't submit a resume and ask "is this good?" We instruct **Gemini 1.5 Pro** to adopt the perspective of a Senior Recruiter, then evaluate the candidate's experience *contextually* against the specific requirements of the live Job Description.

What the Model Evaluates

Contextual Fit

Experience relevance, not just title matching

Justified Score

Numeric rating with a written rationale

Action Recommendation

INTERVIEW / REVIEW /
REJECT with reasoning

Made with **GRAMMA**

Quantifiable Efficiency Gains

PROJECTED ROI

32→15

Days to Hire

Target reduction from 32 days to under 15 days — cutting time-to-offer in half.

20hrs

Weekly Time Saved

Complete elimination of manual resume screening time for Sarah — reclaimed for high-value recruiting work.

100%

Standardized Evaluation

Every applicant scored against the same JD criteria — bias-reduced, consistent, and fully auditable from day one.

📌 **The compounding effect:** Faster screening → earlier offers → higher acceptance rates → reduced cost-per-hire. Each metric improvement reinforces the others.

Made with GRAMMA

Next Steps: Moving from MVP to Production

Current Status

01

Functional MVP Complete

Backend logic fully tested via Postman. All four workflow nodes validated end-to-end.

02

High-Fidelity Frontend Live

React prototype built in Lovable. Upload → Score → Review flow fully functional.

03

Gemini Integration Proven

Structured JSON output confirmed. Scoring accuracy validated on test resume sets.

Phase 2 Ask

We are requesting approval to advance Recruit-AI from MVP to production-ready infrastructure:

- ☐ The workflow is proven. The ask is infrastructure — not experimentation.

1

Vector Database Integration


Pinecone or Supabase for persistent candidate storage, search, and longitudinal comparison across roles.

2

Secure Authentication

Role-based access control to ensure recruiter-level and admin-level permissions are properly enforced.

UX/Wireframes



[link to excalidraw](#)

User Flow Documentation: Recruit-AI

Target User: Sarah (Solo Talent Acquisition Manager) **Primary Goal:** Upload a batch of unstructured resumes and receive a prioritized shortlist of candidates for a specific role.

1. Entry Point: The Dashboard Overview

- **Action:** Sarah logs into the Recruit-AI portal.
- **UI State:** She lands on the main Dashboard, viewing the top-level metric cards (Active Roles, Resumes Screened, Pending Reviews).

- **Friction Removed:** No cluttered navigation; she immediately sees the health of her hiring pipeline.

2. Step 1: Defining the Target (Job Openings)

- **Action:** Sarah navigates to the "Job Openings" tab and clicks "Create New Role." She pastes the Job Title and the rigid Job Description into the modal and saves.
- **System Logic:** This text is temporarily stored in the local state. It will act as the baseline criteria for the AI evaluation in the next step.

3. Step 2: The Bulk Upload (Candidates Page)

- **Action:** Sarah navigates to the "Candidates" tab and drags a batch of newly received PDF resumes into the drag-and-drop zone.
- **System Logic:** The frontend captures the files and triggers a webhook to our n8n backend workflow.


4. Step 3: Agentic Processing State

- **UI State:** The drag-and-drop zone transforms into a loading spinner reading *"Screening Candidates..."* to provide immediate system feedback.
- **Backend Logic (Invisible to User):** 1. The n8n webhook receives the PDF. 2. A data extraction node parses the unstructured text. 3. The text, along with the saved Job Description, is passed to the AI model with a strict prompt to act as an evaluator. 4. The AI outputs a JSON response containing a Match Score, Recommendation, and Summary.

5. Resolution: The Actionable Outcome

- **UI State:** The loading state clears. The Candidates table instantly populates with the ranked results.
- **Decision Point:** Sarah reviews the top-ranked candidate (e.g., "Match Score: 88, INTERVIEW"). She clicks the expandable summary to read the AI's justification.
- **Outcome:** Instead of spending 20 minutes reading, Sarah spends 10 seconds reviewing the data and immediately clicks a button to schedule the interview. The "Resumes Screened" dashboard metric increments by 1.

Recruit-AI User Flow for Batch Resume Screening



Frontend link

[THIS IS THE FRONTEND LOVABLE LINK](#)

n8n Backend flow

[LINK TO n8n WORKFLOW](#)

Metrics and roadmap

Deliverable 6: Metrics & Product Roadmap

Project: Recruit-AI **Product Manager:** [Your Name]

Part 1: Success Metrics (KPIs)

To measure the true impact of Recruit-AI, we are categorizing our telemetry into three standard PM buckets: North Star, Efficiency (Leading), and Guardrail metrics.

1. The North Star Metric

- **Time-to-Shortlist (TTS):** The total time elapsed from the moment a job is posted to the moment the Talent Acquisition team has 10 qualified candidates ready to interview.
 - *Baseline:* 14–21 days (Industry average for high-volume tech roles).
 - *Target:* < 24 hours.

2. Efficiency & Operational Metrics (Leading Indicators)

- **Recruiter Hours Saved per Requisition:** * *Baseline:* ~23 hours spent manually screening 500+ unstructured PDFs.
 - *Target:* < 2 hours (Recruiter only reviews the AI's top 10% shortlist).
- **Cost Per Screen:**
 - *Calculation:* (API Token Cost + Cloud Compute) / Number of Resumes.
 - *Target:* < ₹5 (or \$0.05) per resume screened, making it infinitely cheaper than human hourly rates.
- **System Latency:** Time taken to parse, analyze, and output UI results per document.
 - *Target:* < 4 seconds per PDF via the n8n + Gemini workflow.

3. Guardrail & Quality Metrics (Risk Mitigation)

- **Hiring Manager Acceptance Rate (HMAR):** Of the candidates the AI scored as "INTERVIEW," how many did the human hiring manager actually agree to speak with?
 - *Target:* > 80%. (If this drops, it means our LLM prompt is hallucinating or misaligned with the Job Description).
- **False Negative Audit Rate:** Bi-weekly random sampling of 50 AI-rejected resumes by a human to ensure the AI isn't accidentally filtering out non-traditional, high-quality talent due to strict formatting bias. *Target:* < 2% error rate.



Strategic Product Roadmap (Now, Next, Later)

This roadmap dictates how we evolve Recruit-AI from an internal operational MVP into a deeply integrated, enterprise-ready workflow engine.

Phase 1: NOW (v1.0 - The MVP Validation)

Objective: Prove the core AI evaluation loop and eliminate the manual reading bottleneck.

- **Core Deliverables:**

- Deploy the React/Lovable frontend dashboard for intuitive file drop.
- Implement zero-shot PDF parsing via n8n.
- Integrate Gemini 1.5 Pro to output structured JSON (Match Score out of 100, Action Badge, Justification).
- **Status:** Completed & Ready for testing.

Phase 2: NEXT (v1.5 - Integration & Retention)

Objective: Move from a standalone tool to an embedded system that talks to Sarah's existing tech stack.

- **Core Deliverables:**


- **ATS API Connectors:** Build two-way sync with legacy systems like Greenhouse and Workable. Recruit-AI acts as the AI layer on top of their storage layer.
- **Vector Database (RAG):** Implement Supabase or Pinecone. Instead of evaluating resumes once, we store their embeddings. Sarah can prompt: *"Show me candidates who applied last year who know React."*
- **Automated Candidate Comms:** Trigger immediate webhook emails to reject low-scoring candidates (improving employer brand via rapid response).

Phase 3: LATER (v2.0 - Autonomous Agentic Workflows)

Objective: Complete end-to-end automation of the top-of-funnel pipeline.

- **Core Deliverables:**

- **Conversational AI Pre-Screening:** Deploy a WhatsApp/SMS agent that pings top candidates to verify salary expectations and notice periods *before* Sarah ever talks to them.
- **Auto-Scheduling:** Grant the agent OAuth access to Sarah's Google Calendar to negotiate and book interview slots autonomously.
- **Outbound Sourcing:** Chrome extension integration to score LinkedIn profiles on the fly against open requisitions.



Process and decision log

Process & Decision Log

Project: Recruit-AI **Product Manager:** [Akshay]

1. Executive Summary

The objective of this MVP was to build a functional, end-to-end proof of concept for an Agentic AI screening tool within a highly constrained timeframe. To achieve this, the product strategy heavily favored "Build Speed" and "Visual Fidelity" over "Enterprise Scalability" for version 1.0.

Below is the log of key architectural and product decisions made during Phase 3 (The Build).

2. Key Product & Technical Decisions

Decision 1: Frontend Prototyping (Choosing Lovable)

- **The Goal:** Build a responsive, modern React dashboard that simulates a high-fidelity user experience without writing thousands of lines of boilerplate code.
- **The Alternatives Considered:** Figma (non-functional), standard React/Vite build from scratch (too slow for the MVP timeline), Retool/Bubble (steep learning curve for custom UI styling).
- **The Choice: Lovable.**
- **The Trade-off (Why):** Lovable allowed for rapid, AI-assisted generation of a functional React application. While it relies on mock data for the deployed public URL, it provided the exact interactive UX (drag-and-drop zones, loading states, expanding table rows) required to successfully pitch the workflow to stakeholders. It prioritizes pitch readiness over immediate backend integration.

Decision 2: Backend Orchestration (Choosing n8n)

- **The Goal:** Orchestrate a complex data flow (Webhook → PDF Parsing → LLM Prompting → JSON Output) visually so the logic could be easily audited and adjusted.
- **The Alternatives Considered:** Zapier (too expensive at scale, rigid), custom Node.js/Python microservice (unnecessary technical overhead for an MVP).
- **The Choice: n8n (Local Execution).**
- **The Trade-off (Why):** n8n provides a visual, node-based architecture that is perfect for prototyping agentic workflows. By running it locally (via Cloudflare Tunnels for webhook exposure), we bypassed complex cloud deployment and CORS restrictions. It allowed for rapid iteration of the AI prompt and immediate testing via Postman.

Decision 3: The AI / LLM Engine (Choosing Gemini 2.5 Flash)

- **The Goal:** Accurately parse unstructured, messy PDF text and evaluate it strictly against a supplied Job Description at high speeds.
- **The Alternatives Considered:** Gemini 1.5 Pro, Anthropic Claude 3.5 Sonnet.
- **The Choice: Google Gemini 2.5 Flash.**
- **The Trade-off (Why):** While "Pro" models offer deep reasoning for highly complex tasks, Gemini 2.5 Flash is explicitly optimized for low-latency, high-volume workloads—making it the perfect engine for a bulk resume screening tool. It maintains a massive 1-million-token context window (essential for reading dense, multi-page PDFs) while significantly reducing API costs and processing time. Additionally, its native support for Structured Outputs guaranteed that it would consistently return the clean JSON (Match Score, Action Badge, Justification) required to keep our frontend UI from breaking

Decision 4: The B2B MVP "CORS" Workaround

- **The Roadblock:** Connecting a live browser frontend directly to a local n8n instance resulted in immediate Cross-Origin Resource Sharing (CORS) security blockers.
- **The PM Decision:** Instead of spending days configuring reverse proxies and server headers—which distracts from the core product value proposition—we decoupled the demo.
- **The Execution:** We built the backend logic to be proven via API tools (Postman) and hardcoded the frontend state to simulate the exact 3-second processing delay and output. This ensured the CTO pitch could proceed flawlessly without getting bogged down in local networking bugs.

3. Key Learnings & Next Steps

- **Prompt Engineering is the New Backend:** The hardest part of the build was not the routing, but refining the LLM prompt to stop "hallucinating" candidate skills. We learned that the AI must be explicitly instructed to *only* use the provided PDF text and to adopt the persona of a strict, senior technical recruiter.
- **Next Step for Production:** For Phase 2, the immediate technical requirement is migrating the n8n workflow to a managed cloud instance and implementing a proper backend server (e.g., Node/Express) to handle CORS securely between the Lovable frontend and the AI endpoints.

Video pitch

[pitch video link](#)